



CREATING RELIABLE, DISTRIBUTED APPLICATIONS

CLOUD HASKELL

```
module AboutMe where
import Data.Human
```

```
theSpeaker :: Human
```

```
theSpeaker = Human
```

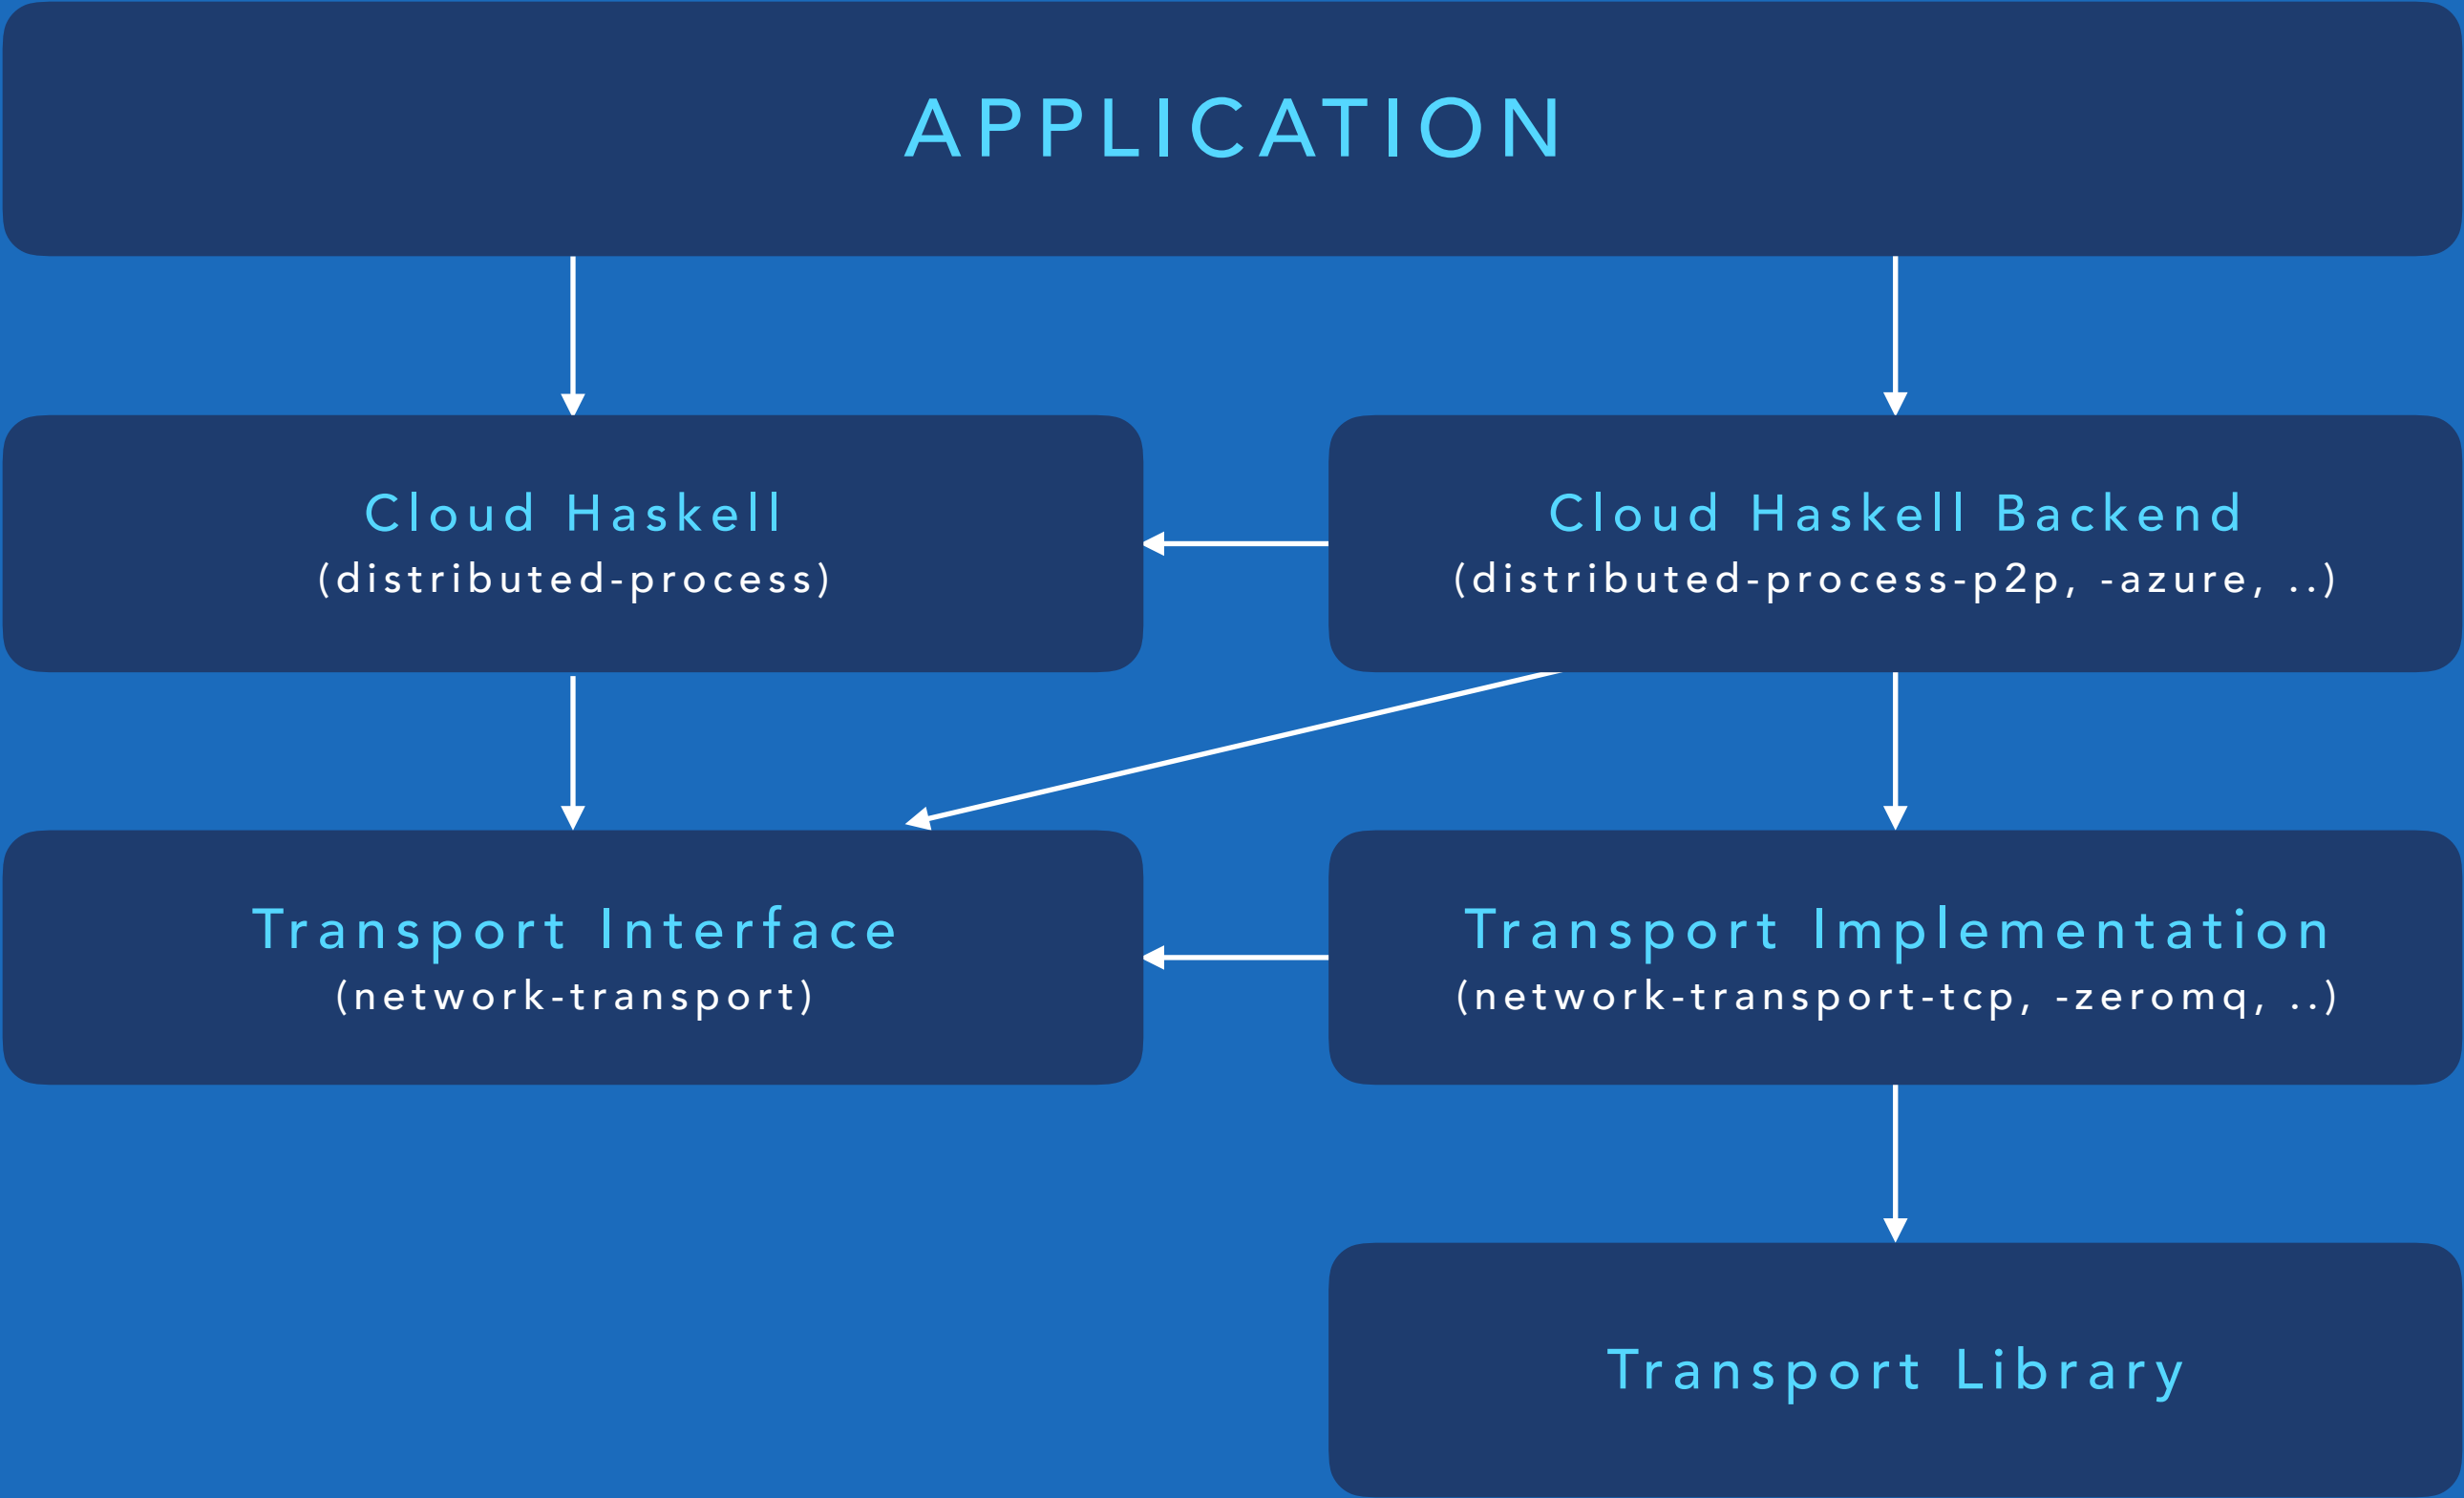
```
    { names    = [ mkName "sn" "不动点帕琪"
                    , mkName "en" "Tyler Ling"
                    , mkName "zh" "凌辉" ]
    , email    = "tylerblugandersen@gmail.com"
    , qq       = "503228590"
    , langs    = [ "Chinese", "English"
                    , "Haskell", "C++", "Lua" ]
    }
```

WHY HASKELL

- 纯函数：引用透明，不可变量
- 强类型：保障代码和被传递数据的正确性
- Monad：表达、控制副作用

分布式编程的挑战

- 可扩展性需求
 - 通讯后端多样 (TCP/IP, UDP, ZeroMQ, In-memory, Pipeline...)
 - 节点种类丰富 (物理机器, 虚拟机, 云服务器...)
 - 组织形式灵活 (Master-slave, P2P, ...)
- 节点间的交互
 - 消息的发送与接收
 - 启动远程进程
- 容错机制
 - 错误检测
 - 错误恢复

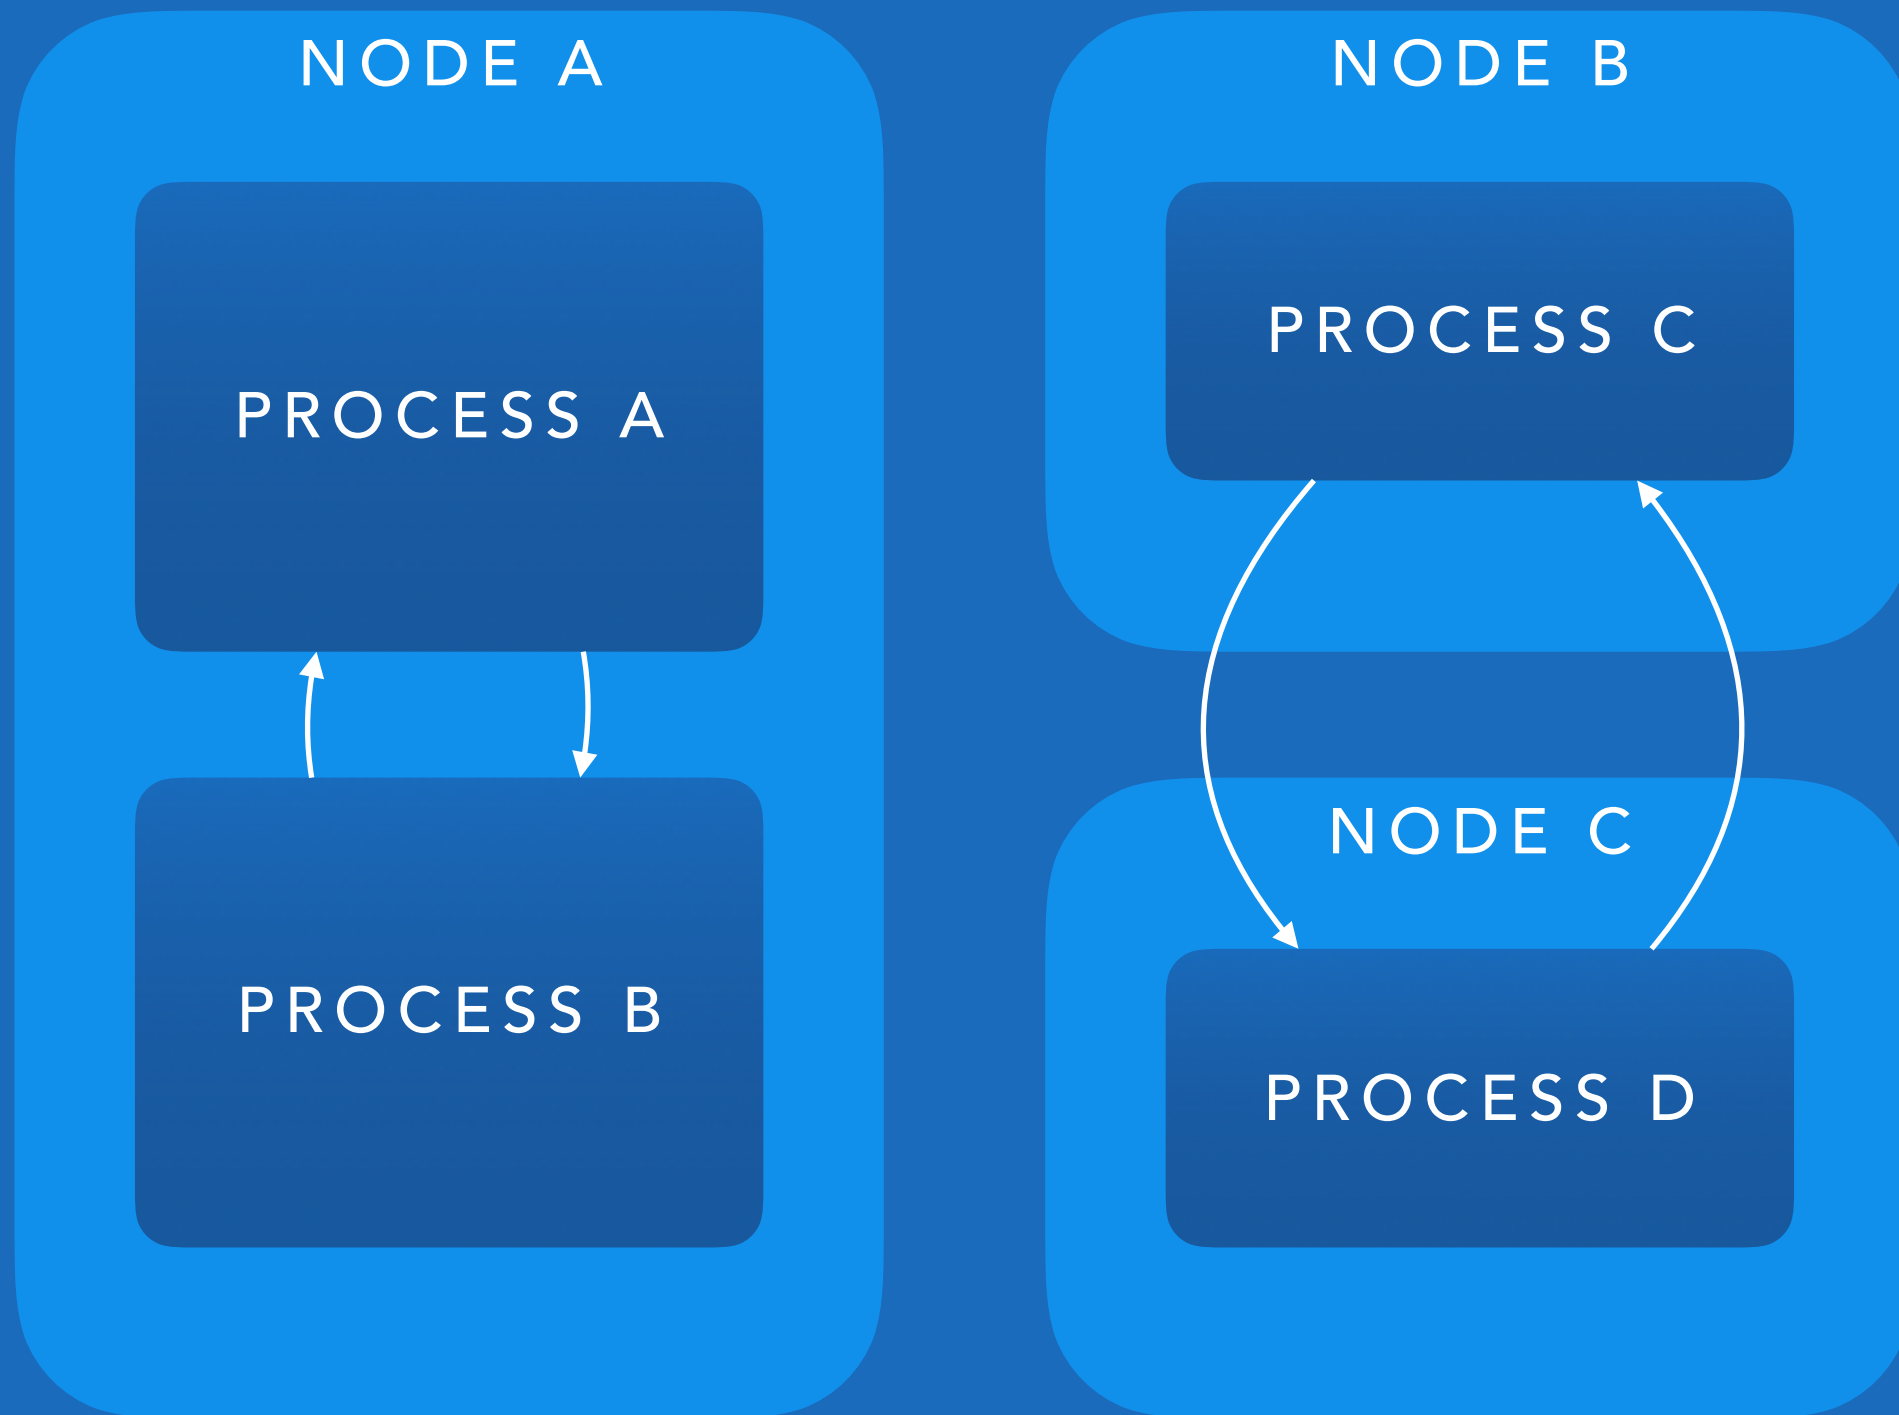


“Cloud Haskell is a set of libraries that bring Erlang-style concurrency and distribution to Haskell programs.”

CLOUD HASKELL/DOCUMENTATION

抽象

Actor Model



抽象

- Node
 - `newLocalNode :: Transport -> RemoteTable -> IO LocalNode`
 - `forkProcess :: LocalNode -> Process () -> IO ProcessId`
 - `runProcess :: LocalNode -> Process () -> IO ()`
- Process
 - 轻量级
 - 无共享资源
 - 异步消息传递 (send & receive)
 - 启动本地、远程进程 (spawnLocal & spawn)
 -

PROCESS LAYER

- Message
 - 有限的可序列化数据结构。
 - 类型一致性检查。

```
class (Binary a, Typeable a) => Serializable a
```

- | | |
|---------------------------------|--------------------|
| • Serializable | • Non-serializable |
| • Int, Char, etc. | • MVar, TVar, etc. |
| • (Serializable a) => [a], etc. | • LocalNode, etc. |

PROCESS LAYER

- Process
 - It's **Monad!** (also **MonadIO**)
 - ProcessId: 用于接收数据的唯一的身份标识符。

```
send :: Serializable a => ProcessId -> a -> Process ()
```

```
expect :: Serializable a => Process a
```

- | | |
|---|---|
| <ul style="list-style-type: none">• send<ul style="list-style-type: none">• 异步• 绝不失败：但不保证接收端能否收到 | <ul style="list-style-type: none">• expect<ul style="list-style-type: none">• 同步：没有收到预期类型的消息前堵塞线程• 其他不符合的消息将继续留在消息信箱中。 |
|---|---|

PROCESS LAYER

```
data Ping = Ping ProcessId
    deriving (Typeable, Generic, Binary)

data Pong = Pong ProcessId
    deriving (Typeable, Generic, Binary)
```

```
ping :: Process ()
ping = do
    self <- getSelfPid
    Pong pid <- expect
    send pid (Ping self)
    ping
```

```
{-# LANGUAGE DeriveDataTypeable #-}

{-# LANGUAGE DeriveGeneric #-}

{-# LANGUAGE DeriveAnyClass #-}
```

PROCESS LAYER

- Multiple Choices

```
data Ping = Ping ProcessId
```

```
data Pong = Pong ProcessId
```

```
pingAndPong :: Process ()
```

```
pingAndPong = do
```

```
    self <- getSelfPid
```

```
    ..?
```

PROCESS LAYER

- Multiple Choices

```
pingAndPong :: Process ()
```

```
pingAndPong = do
```

```
  self <- getSelfPid
```

```
  receiveWait
```

```
    Ping pid -> send pid (Pong self)
```

```
    Pong pid -> send pid (Ping self)
```

PROCESS LAYER

- Multiple Choices

- `match :: Serializable a => (a -> Process b) -> Match b`
- `receiveWait :: [Match b] -> Process b`

```
pingAndPong :: Process ()
```

```
pingAndPong = do
```

```
  self <- getSelfPid
```

```
  receiveWait
```

```
    [ match $ \(Ping pid) -> send pid (Pong self)
```

```
    , match $ \(Pong pid) -> send pid (Pong self) ]
```

PROCESS LAYER

- Multiple Choices

- `match :: Serializable a => (a -> Process b) -> Match b`
- `receiveWait :: [Match b] -> Process b`
- `receiveTimeout :: Int -> [Match b] -> Process (Maybe b)`
- `matchIf :: Serializable a => (a -> Bool) -> (a -> Process b) -> Match b`
- `matchAny :: (Message -> Process b) -> Match b`
-

`expect :: Serializable a => Process a`

`expect = receiveWait [match return]`

PROCESS LAYER

- What if...

```
Process 1> send pid2 $ length [1, 2, 3]
```

```
Process 2> v <- except  
            send pid3 (v + 2 :: Double)
```

```
Process 3> Excuse me?
```

“Let there be type!”

FIXED-POINT PATCHY

PROCESS LAYER

- Typed Channel
 - SendPort (Serializable): 发送端。
 - ReceivePort (Non-serializable): 接收端。

```
newChan :: Serializable a => Process (SendPort a, ReceivePort a)
```

```
sendChan :: Serializable a => SendPort a -> a -> Process ()
```

```
receiveChan :: Serializable a => ReceivePort a -> Process a
```

```
receiveChanTimeout :: Serializable a => Int -> ReceivePort a -> Process (Maybe a)
```

```
mergePortsBiased :: Serializable a => [ReceivePort a] -> Process (ReceivePort a)
```

```
mergePortsRR :: Serializable a => [ReceivePort a] -> Process (ReceivePort a)
```

PROCESS LAYER

```
example :: Process ()
```

```
example = do
```

```
    (sp, rp) <- newChan
```

```
    -- spawnLocal :: Process () -> Process ProcessId
```

```
    spawnLocal $ sendChan sp "Hello world"
```

```
    "Hello world" <- receiveChan rp
```

But...

spawnLocal

PROCESS LAYER

- 创建远程进程

```
sender :: SendPort String -> Process ()
```

```
sender sp = sendChan sp "Hello world"
```

```
example :: NodeId -> Process ()
```

```
example nid = do
```

```
    (sp, rp) <- newChan
```

```
    spawn ???
```

PROCESS LAYER

- 创建远程进程

`spawn :: NodeId -> Closure (Process ()) -> Process ProcessId`

Where to run

What to run

PROCESS LAYER

- 创建远程进程
 - What might be Closure?
 - 函数本体
 - 调用环境（自由变量，引用库.....）

```
data Closure a = Closure function environment
```


PROCESS LAYER

- 创建远程进程
 - What might be Closure?
 - 函数本体
 - 调用环境（自由变量，引用库.....）
 - 这样的 Closure 能被序列化吗？

```
instance Binary (a -> b) where
```

```
  put x = ____ ???
```

```
  get   = ____ ???
```

PROCESS LAYER

- 创建远程进程
 - What might be Closure?
 - 函数本体
 - 调用环境（自由变量，引用库.....）
 - 这样的 Closure 能被序列化吗？
 - × 函数本体
 - × 调用环境

PROCESS LAYER

- 创建远程进程
 - 解决方案
 - 扩展运行时
 - 致使代码失去对序列化的精准控制。
 - 误将序列化无意义的数据与函数一同传至其他节点。
 - 一些数据经过人为处理后传输更有效率。
 - **monolithic + static value**

PROCESS LAYER

- 创建远程进程
 - 哪些函数最容易在其他节点上引用？
 - 其他节点也存在同样的代码。
 - 顶级 (Top-level) 。
 - 无自由变量，或者自由变量也是顶级的。

```
instance Serializable (Static a)
```

```
static :: 对于满足条件的 a => a -> Static a
```

```
unstatic :: Static a -> a
```

PROCESS LAYER

- 创建远程进程
 - 实现 Static value
 - RemoteTable (Map String Dynamic)
 - Since GHC 7.10: GHC.StaticPtr

```
newtype Static a = Static StaticLabel
data StaticLabel =
    StaticLabel String
  | StaticApply !StaticLabel !StaticLabel
#if __GLASGOW_HASKELL__ >= 710
  | StaticPtr SDynamic
    -- data SDynamic = SDynamic TypeRep (StaticPtr GHC.Any)
#endif
```

PROCESS LAYER

data Closure a where

Closure :: Serializable env => Static (env -> a) -> env -> Closure a

instance Binary (Closure a) where

put (Closure f env) = put f >> put env

get = ????



Which deserializer to use?

PROCESS LAYER

```
data Closure a = Closure (Static (ByteString -> a)) ByteString
```

Deserializer



Environment

```
unclosure :: Typeable a => RemoteTable -> Closure a -> Either String a
unclosure rtable (Closure dec env) = do
  f <- unstatic rtable dec
  return (f env)
```


PROCESS LAYER

```
sender :: SendPort String -> Process ()  
sender sp = sendChan sp "Hello world"
```

```
senderStatic :: Static (SendPort String -> Process ())  
senderStatic = staticLabel "$sender"
```

```
decodeSendPortStatic :: Static (ByteString -> SendPort String)  
decodeSendPortStatic = staticLabel "$decodeSendPort"
```

```
senderClosure :: SendPort String -> Closure (Process ())  
senderClosure sp = closure decoder (encode sp)  
  where decoder :: Static (ByteString -> Process ())  
        decoder = senderStatic `staticCompose` decodeSendPortStatic
```

PROCESS LAYER

```
rtable :: RemoteTable
rtable =
    registerStatic "$sender" (toDynamic sender)
  . registerStatic "$decodeSendPort"
    (toDynamic (decode :: ByteString -> SendPort String))
  $ initRemoteTable
```

```
newLocalNode :: Transport -> RemoteTable -> IO LocalNode
```

“Template Haskell!”

FIXED-POINT PATCHY

PROCESS LAYER

```
{-# LANGUAGE TemplateHaskell #-}

import Control.Distributed.Process
import Control.Distributed.Process.Closure
import Control.Distributed.Process.Node
import Network.Transport.TCP (createTransport, defaultTCPParameters)

sender :: SendPort String -> Process ()
sender sp = sendChan sp "Hello world"

remotable ['sender]
-- __remoteTable :: RemoteTable -> RemoteTable

remoteTable :: RemoteTable
remoteTable = Main.__remoteTable initRemoteTable

main :: IO ()
main = do
    Right transport <- createTransport "127.0.0.1" "10001" defaultTCPParameters
    node <- newLocalNode transport remoteTable
    runProcess node $ do
        snid <- getSelfNode
        (sp, rp) <- newChan
        spawn snid ($\mkClosure 'sender) sp)
```

PROCESS LAYER

- 创建远程进程
 - Polymorphic?
 - Data.Rank1Dynamic (toDynamic)
 - Data.Rank1Typeable (Typeable, ANY, ANY1, ANY2, ANY3, ANY4)

```
rtable :: RemoteTable
rtable =
  registerStatic "$decode"
    (toDynamic (decode :: ByteString -> ANY)) -- Really?
  $ initRemoteTable
```

PROCESS LAYER

```
data SerializableDict a where
  SerializableDict :: Serializable a => SerializableDict a
  deriving (Typeable)

decodeDict :: SerializableDict a -> ByteString -> a
decodeDict SerializableDict = decode

rtable :: RemoteTable
rtable =
  registerStatic "$decodeDict"
    (toDynamic (decodeDict :: SerializableDict ANY -> ByteString -> ANY))
  $ initRemoteTable

staticDecode :: Typeable a => Static (SerializableDict a) -> Static (ByteString -> a)
staticDecode dict = decodeDictStatic `staticApply` dict
  where
    decodeDictStatic :: Typeable a => Static (SerializableDict a -> ByteString -> a)
    decodeDictStatic = staticLabel "$decodeDict"
```

PROCESS LAYER

```
sdictT :: SerializableDict T
sdictT = SerializableDict
$(mkStatic 'sdictT) :: Static (SerializableDict T)
```

```
remotable ['f] -- f :: T1 -> T2
$(functionSDict 'f) :: Static (SerializableDict T1)
-- if f :: T1 -> Process T2
$(functionTDict 'f) :: Static (SerializableDict T2)
```


PROCESS LAYER

- See also...
 - `Control.Distributed.Process.Closure` (distributed-process)

```
type CP a b = Closure (a -> Process b)
```

```
idCP :: Typeable a => CP a a
```

```
returnCP :: Serializable a => Static (SerializableDict a) -> a -> Closure (Process a)
```

```
bindCP :: (Typeable a, Typeable b) => Closure (Process a) -> CP a b -> Closure (Process b)
```

```
seqCP :: (Typeable a, Typeable b) => Closure (Process a) -> Closure (Process b) -> Closure (Process b)
```

```
cpLink :: ProcessId -> Closure (Process ())
```

```
cpSend :: Typeable a => Static (SerializableDict a) -> ProcessId -> CP a ()
```

```
cpExpect :: Typeable a => Static (SerializableDict a) -> Closure (Process a)
```

PROCESS LAYER

- 容错
 - Link
 - 单向: (A) Process -> (B) Process / Node / Channel
 - 当 B 正常/异常结束、或失去联系时, A 会产生一个异步异常 (ProcessLinkException), 导致 A 也被终止。
 - ProcessLinkException 没有被 Cloud Haskell 导出。

```
data ProcessLinkException = ProcessLinkException ProcessId DiedReason
```

```
link :: ProcessId -> Process ()
```

```
unlink :: ProcessId -> Process ()
```

```
-- located in distributed-process-extras, implemented by moniter
```

```
linkOnFailure :: ProcessId -> Process ()
```

PROCESS LAYER

- 容错
 - Monitor
 - 单向
 - 当被监控的 Process/Node/Port 结束时，进行监控的进程将会收到类型为 Process-/Node-/PortMonitorNotification 的消息。
 - 每次对 monitor 的调用产生新的 MonitorRef，需要分别 unmonitor。

```
data ProcessMonitorNotification =  
    ProcessMonitorNotification MonitorRef ProcessId DiedReason  
  
monitor :: ProcessId -> Process MonitorRef  
  
unmonitor :: MonitorRef -> ProcessId  
  
withMonitor :: ProcessId -> Process a -> Process a
```

PROCESS LAYER

```
linkOnFailure :: ProcessId -> Process ()
linkOnFailure them = do
  us <- getSelfPid
  tid <- liftIO $ myThreadId
  void $ spawnLocal $ do
    callerRef <- P.monitor us
    calleeRef <- P.monitor them
    reason <- receiveWait [
      matchIf (\(ProcessMonitorNotification mRef _ _) ->
        mRef == callerRef) -- nothing left to do
        (\_ -> return DiedNormal)
      , matchIf (\(ProcessMonitorNotification mRef' _ _) ->
        mRef' == calleeRef)
        (\(ProcessMonitorNotification _ _ r') -> return r')
    ]
  case reason of
    DiedNormal -> return ()
    _ -> liftIO $ throwTo tid (ProcessLinkException us reason)
```

PROCESS LAYER

```
spawnLink :: NodeId -> Closure (Process ()) -> Process ProcessId
```

```
spawnMonitor :: NodeId -> Closure (Process ()) -> Process (ProcessId, MonitorRef)
```

```
spawnSupervised :: NodeId -> Closure (Process ()) -> Process (ProcessId, MonitorRef)
```

PROCESS LAYER

- See also...
 - `Control.Distributed.Process.Supervisor` (distributed-process-supervisor)
 - Supervision tree: Hierarchical process structure.
 - Restart Strategies...

What's the behavior of ...

link pid	<- Async!
send pid "Hello world"	<- Async!
unlink pid	<- Async!

```
link pid
send pid "Hello world"
reply <- expect
unlink pid
```

PROCESS LAYER

- Process Layer 的特点： 底层
 - 类型安全保障弱。
 - 手动错误侦测与恢复。
 - 对逻辑的表达能力弱。

TASKレヤの消失

- 在很久很久以前..... (Cloud Haskell 只有一个叫 remote 的库的时候)
 - Promise/future (inspired by Skywriting/CIEL)
 - 代表一个完成或没有完成的计算结果, Serializable。
 - 对值的提取操作会让该 Promise 在得到结果前堵塞。
 - 计算在 TaskM 中完成。
 - TaskM: 不能执行任何 IO 操作的 Monad
 - 最大限度减少错误的可能。
 - 一个出错终止的 TaskM 可以很容易地自动重启。

```
newPromise :: Serializable a => Closure (TaskM a) -> TaskM (Promise a)
readPromise :: Serializable a => Promise a -> TaskM a
runTask :: Serializable => TaskM a -> Process a
```

TASKレヤの消失

```
avg :: [Integer] -> TaskM Integer
avg xs = return $ sum xs `div` fromIntegral (length xs)

diff :: Promise Integer -> Promise Integer -> TaskM Integer
diff pa pb = do
  a <- readPromise pa
  b <- readPromise pb
  return $ (a + b) / 2

$(remotable ['avg, 'diff])

process :: Process ()
process = do
  res <- runTask $ do
    p1 <- newPromise ($ (mkClosure 'avg) [0..50])
    p2 <- newPromise ($ (mkClosure 'avg) [50..100])
    p3 <- newPromise ($ (mkClosure 'diff) p1 p2)
    readPromise p3
  say $ "Result: " ++ show res
```

TASKレヤの消失

- 消失の原因
 - 不可扩展：Master-slave
 - 不够健壮：Master 节点一旦崩溃整个集群必须重启。
 - 节点分配的算法实用性低：round-robin
 - 需要硬件资源监控和负载均衡的技术。

Maybe it will come back...

TASKレヤの消失

- Control.Distributed.Process.Async!
 - 分布式版本的 Control.Concurrent.Async。
 - 同时支持堵塞等待与非堵塞式查询。
 - 不限制 IO 的使用。

```
async :: Serializable a => AsyncTask a -> Process (Async a)
```

```
asyncLinked :: Serializable a => AsyncTask a -> Process (Async a)
```

```
wait :: Async a -> Process (AsyncResult a)
```

```
poll :: Serializable a => Async a -> Process (AsyncResult a)
```

TASKレヤの消失

```
data AsyncTask a =  
    AsyncTask {  
        asyncTask :: Process a  
    }  
| AsyncRemoteTask {  
    asyncTaskDict :: Static (SerializableDict a)  
    , asyncTaskNode :: NodeId  
    , asyncTaskProc :: Closure (Process a)  
    }  
  
task :: Process a -> AsyncTask a  
remoteTask :: Static (SerializableDict a) -> NodeId -> Closure (Process a) -> AsyncTask a
```

MANAGED PROCESS

- Client/Server 模式 (distributed-process-client-server)
- 自动消息解码、分发以及错误处理
 - Mailbox -> 用户定义的 handlers (根据消息的类型以及用户提供的 predicates)
- 两种可选的通讯方法
 - cast: 客户端异步发送消息, 服务端不回复。
 - call: 远程过程调用 (Remote Procedure Call), 客户端等待服务端传回结果。

```
cast :: (Addressable a, Serializable m) => a -> m -> Process ()
```

```
call :: (Addressable s, Serializable a, Serializable b) => s -> a -> Process b
```

MANAGED PROCESS

```
import Data.Time.LocalTime

data TimeType = TtUTC | TtLocal
  deriving (Eq, Typeable, Generic, Binary)
data GetTime = GetTime TimeType
  deriving (Typeable, Generic, Binary)

getTime :: (Addressable a) => a -> TimeType -> Process String
getTime a tt = call pid $ GetTime tt

timeServer :: Process ProcessId
timeServer =
  let server = statelessProcess {
    apiHandlers =
      [ handleCallIf_ (input $ \(GetTime tt) -> tt == TtUTC)
        (\_ -> fmap show (liftIO getZonedTime))
      , handleCallIf_ (input $ \(GetTime tt) -> tt == TtLocal)
        (\_ -> fmap (show . zonedTimeToLocalTime) $ liftIO getZonedTime)]
    , unhandledMessagePolicy = Drop }
  in spawnLocal serve () (statelessInit Infinity) server
```

参考

- Functional programming for the data centre, June 2011
 - Jeffrey Epstein
- Towards Haskell in the Cloud, September 2011
 - Jeff Epstein
 - Andrew P. Black
 - Simon Peyton-Jones
- Cloud Haskell Documentation, Tutorials
 - <http://haskell-distributed.github.io>

THANKS!